

**SIMATS ENGINEERING
ASSESSMENT TOOL 3**

COURSE NAME: Fundamentals of Data Science

COURSE CODE: DSA04

COURSE OUTCOME COVERED

***CO1:** Analyse the role of data science, facets of data, and the big data ecosystem for informed decision-making. (BL2)*

QUESTIONS: 05

Total Marks:100

Annexure A

Exploratory Data Analysis - Based Practical Assessment

Code of Conduct:

*I **A.Sai Rohit (192472144)** (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.
I understand that any violation of academic integrity rules will result in disciplinary action.*

Signature of the student: A.Sai Rohit

Annexure A

Exploratory Data Analysis-Based Practical Assessment

1. Education Analytics Dataset

Question 1

A college wants to analyze student performance based on attendance, study hours, internal marks, and final result. Perform EDA to identify the factors affecting student results.

Student_ID	Attendance	Study_Hours	Internal_Marks	Result
S01	92	5	45	Pass
S02	55	2	24	Fail
S03	78	4	38	Pass
S04	60	3	28	Fail
S05	88	5	42	Pass
S06	47	1	20	Fail
S07	95	6	48	Pass
S08	72	3	35	Pass
S09	50	2	22	Fail
S10	83	4	40	Pass
S11	66	3	30	Fail
S12	90	5	44	Pass
S13	58	2	25	Fail
S14	76	4	36	Pass
S15	62	3	29	Fail
S16	85	5	41	Pass
S17	45	1	18	Fail
S18	80	4	39	Pass
S19	68	3	32	Pass
S20	53	2	23	Fail

2. Healthcare Analytics Dataset

Question 2

A hospital wants to analyze patient health risk based on age, blood sugar, blood pressure, BMI, and disease status. Perform EDA to find important health patterns.

Patient_ID	Age	Sugar_Level	BP	BMI	Disease_Status
P01	45	145	130	28	Yes
P02	32	98	118	23	No
P03	55	180	145	31	Yes
P04	28	90	110	22	No
P05	60	200	150	34	Yes
P06	40	120	125	26	No
P07	52	165	140	30	Yes
P08	35	105	120	24	No
P09	48	155	135	29	Yes
P10	30	95	115	23	No
P11	62	210	155	35	Yes
P12	38	115	122	25	No
P13	50	170	142	30	Yes
P14	27	88	108	21	No

3. Retail Dataset Question 3

A wants to product

on category, price, quantity sold, discount, and revenue. Perform EDA to identify high-selling products and revenue patterns.

Sales

supermarket
analyze
sales based

P15	58	190	148	33	Yes
P16	42	130	128	27	No
P17	46	150	132	28	Yes
P18	34	100	116	24	No
P19	53	175	144	31	Yes
P20	29	92	112	22	No

Product_ID	Category	Price	Quantity_Sold	Discount	Revenue
PR01	Grocery	50	20	5	1000
PR02	Snacks	30	35	3	1050
PR03	Dairy	45	18	2	810
PR04	Fruits	60	25	5	1500
PR05	Vegetables	40	30	4	1200
PR06	Grocery	80	15	6	1200
PR07	Snacks	25	40	3	1000
PR08	Dairy	55	16	2	880
PR09	Fruits	70	22	5	1540
PR10	Vegetables	35	28	4	980
PR11	Grocery	90	12	6	1080
PR12	Snacks	20	45	2	900
PR13	Dairy	65	14	3	910
PR14	Fruits	75	20	5	1500
PR15	Vegetables	45	26	4	1170
PR16	Grocery	100	10	7	1000
PR17	Snacks	35	32	3	1120
PR18	Dairy	50	19	2	950
PR19	Fruits	80	18	6	1440
PR20	Vegetables	30	35	3	1050

4. Banking Loan Approval Dataset

Question 4

A bank wants to analyze loan approval decisions based on income, credit score, loan amount, employment type, and approval status. Perform EDA to identify patterns related to loan approval.

Customer_ID	Income	Credit_Score	Loan_Amount	Employment_Type	Loan_Status
C01	45000	720	200000	Salaried	Approved
C02	25000	580	150000	Self-employed	Rejected
C03	60000	750	300000	Salaried	Approved
C04	22000	560	120000	Unemployed	Rejected
C05	50000	710	250000	Salaried	Approved
C06	30000	600	180000	Self-employed	Rejected
C07	70000	780	350000	Salaried	Approved
C08	28000	590	160000	Self-employed	Rejected
C09	55000	730	270000	Salaried	Approved
C10	24000	570	130000	Unemployed	Rejected
C11	65000	760	320000	Salaried	Approved
C12	35000	620	200000	Self-employed	Rejected
C13	48000	700	240000	Salaried	Approved

C14	26000	585	140000	Unemployed	Rejected
C15	75000	790	400000	Salaried	Approved
C16	32000	610	190000	Self-employed	Rejected
C17	52000	725	260000	Salaried	Approved
C18	27000	575	150000	Self-employed	Rejected
C19	68000	770	330000	Salaried	Approved
C20	23000	555	125000	Unemployed	Rejected

5. Agriculture Crop Yield Dataset

Question 5

An agricultural department wants to analyze crop yield based on rainfall, temperature, fertilizer usage, soil type, and crop yield. Perform EDA to identify the factors influencing crop production.

Farm_ID	Rainfall_mm	Temperature	Fertilizer_kg	Soil_Type	Crop_Yield_kg
F01	850	28	60	Loamy	3200
F02	600	32	45	Sandy	2100
F03	900	27	65	Loamy	3500
F04	500	34	40	Sandy	1800
F05	780	29	55	Clay	3000
F06	650	31	48	Sandy	2300
F07	920	26	70	Loamy	3700
F08	720	30	52	Clay	2800
F09	550	33	42	Sandy	1900
F10	800	28	58	Loamy	3100
F11	880	27	62	Loamy	3400
F12	620	32	46	Sandy	2200
F13	760	29	54	Clay	2950
F14	940	26	72	Loamy	3800
F15	580	33	44	Sandy	2000
F16	820	28	60	Clay	3150
F17	700	30	50	Clay	2700
F18	960	25	75	Loamy	3900
F19	540	34	41	Sandy	1850
F20	790	29	56	Clay	3050

Common EDA Tasks:

1. Identify the number of rows and columns.
2. Check data types of each column.
3. Find missing values and duplicate values.
4. Calculate mean, median, minimum, maximum, and standard deviation.
5. Draw suitable charts such as bar chart, histogram, scatter plot, and box plot.
6. Identify important patterns from the dataset.
7. Write 4 to 5 observations based on EDA.
8. Give a short conclusion from the analysis.

RUBRICS FOR CALCULATION

Exploratory Data Analysis-Based Practical Assessment
--

Criteria	Weightage	Level 4 – Exemplary	Level 3 – Proficient	Level 2 – Developing	Level 1 – Beginning
Dataset Understanding	20	Clearly explains dataset source, features, target variables, data types, and problem context with strong clarity	Explains dataset, features, and variables with adequate clarity	Basic explanation of dataset with limited understanding of variables	Poor understanding of dataset; variables and problem context are unclear
Data Cleaning	25	Effectively identifies and handles missing values, duplicates, incorrect entries, and outliers using suitable methods	Handles most missing values, duplicates, and errors correctly	Performs basic cleaning but misses some important data quality issues	Minimal or incorrect data cleaning; major errors remain in the dataset
Data Analysis and Statistical Summary	20	Provides detailed descriptive statistics, comparisons, patterns, and meaningful interpretations	Provides relevant summary statistics and basic interpretation	Provides limited statistics with weak explanation	Lacks statistical analysis or gives incorrect interpretation
Data Visualization	20	Uses highly appropriate charts with proper labels, titles, and clear visual interpretation	Uses suitable charts with mostly clear labels and explanation	Uses limited or basic charts with unclear interpretation	Poor or incorrect visualizations; charts do not support analysis
Insight Generation and Reporting	15	Presents strong insights, conclusions, and recommendations based on data evidence	Presents clear findings and conclusions with reasonable support	Provides basic findings but lacks depth or proper justification	Findings are unclear, unsupported, or not related to the data

CO1-ASSESSMENT TOOL -2 Concept Map

QUESTIONS : 5

MARKS : 50

Q No	Dataset Understanding (20%)	Data Cleaning (25%)	Data Analysis and Statistical Summary (20%)	Data Visualization (20%)	Insight Generation and Reporting (15%)	Total Score (100%)
1	/2.5	/2.5	/2.0	/2.0	/1.0	/10
2	/2.5	/2.5	/2.0	/2.0	/1.0	/10
3	/2.5	/2.5	/2.0	/2.0	/1.0	/10
4	/2.5	/2.5	/2.0	/2.0	/1.0	/10
5	/2.5	/2.5	/2.0	/2.0	/1.0	/10
Overall Rubrics Value	/12.5	/12.5	/10	/10	/10	/50
	/25	/25	/20	/20	/10	/100

```

In [18]: # Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Create Dataset
data = {
    "Student_ID": ["S01", "S02", "S03", "S04", "S05", "S06", "S07", "S08", "S09", "S10",
                  "S11", "S12", "S13", "S14", "S15", "S16", "S17", "S18", "S19", "S20"],

    "Attendance": [92, 55, 78, 60, 88, 47, 95, 72, 50, 83,
                  66, 90, 58, 76, 62, 85, 45, 80, 68, 53],

    "Study_Hours": [5, 2, 4, 3, 5, 1, 6, 3, 2, 4,
                  3, 5, 2, 4, 3, 5, 1, 4, 3, 2],

    "Internal_Marks": [45, 24, 38, 28, 42, 20, 48, 35, 22, 40,
                     30, 44, 25, 36, 29, 41, 18, 39, 32, 23],

    "Result": ["Pass", "Fail", "Pass", "Fail", "Pass", "Fail", "Pass", "Pass", "Fail", "Fail",
              "Fail", "Pass", "Fail", "Pass", "Fail", "Pass", "Fail", "Pass", "Pass", "Pass"]
}

df = pd.DataFrame(data)

# Display Dataset
print("Education Analytics Dataset")
print(df)

# 1. Number of Rows and Columns
print("\nRows and Columns:", df.shape)

# 2. Data Types
print("\nData Types")
print(df.dtypes)

# 3. Missing Values
print("\nMissing Values")
print(df.isnull().sum())

# Duplicate Values
print("\nDuplicate Values:", df.duplicated().sum())

# 4. Statistical Measures
print("\nMean")
print(df.mean(numeric_only=True))

print("\nMedian")
print(df.median(numeric_only=True))

print("\nMinimum")
print(df.min(numeric_only=True))

```

```

print("\nMaximum")
print(df.max(numeric_only=True))

print("\nStandard Deviation")
print(df.std(numeric_only=True))

# ----- Charts -----

# Histograms for Attendance, Study_Hours, Internal_Marks
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
sns.histplot(df["Attendance"], bins=5, kde=True, ax=axes[0])
axes[0].set_title("Attendance Distribution")
sns.histplot(df["Study_Hours"], bins=6, kde=True, ax=axes[1])
axes[1].set_title("Study Hours Distribution")
sns.histplot(df["Internal_Marks"], bins=5, kde=True, ax=axes[2])
axes[2].set_title("Internal Marks Distribution")
plt.tight_layout()
plt.show()

# Countplot for Result and Boxplot for Attendance vs Result
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.countplot(x="Result", data=df, ax=axes[0])
axes[0].set_title("Pass vs Fail")
sns.boxplot(x="Result", y="Attendance", data=df, ax=axes[1])
axes[1].set_title("Attendance vs Result")
plt.tight_layout()
plt.show()

# Boxplots for Study_Hours vs Result and Internal_Marks vs Result
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.boxplot(x="Result", y="Study_Hours", data=df, ax=axes[0])
axes[0].set_title("Study Hours vs Result")
sns.boxplot(x="Result", y="Internal_Marks", data=df, ax=axes[1])
axes[1].set_title("Internal Marks vs Result")
plt.tight_layout()
plt.show()

# Scatter Plot for Study_Hours vs Internal_Marks and Correlation Heatmap
fig, axes = plt.subplots(1, 2, figsize=(14, 6)) # Increased figsize for better
sns.scatterplot(x="Study_Hours", y="Internal_Marks", hue="Result", data=df, ax=
axes[0].set_title("Study Hours vs Internal Marks")

numeric_df = df[["Attendance", "Study_Hours", "Internal_Marks"]]
sns.heatmap(numeric_df.corr(), annot=True, cmap="Blues", ax=axes[1])
axes[1].set_title("Correlation Matrix")
plt.tight_layout()
plt.show()

# 5. Important Patterns
print("\nImportant Patterns")
print("• Students with higher attendance mostly passed.")
print("• Students who studied more hours scored higher marks.")
print("• Higher internal marks are associated with passing.")

```



```
print("• Attendance and Internal Marks have a positive correlation.")

# 6. Observations
print("\nObservations")
print("1. Higher attendance improves the chances of passing.")
print("2. Students studying 4-6 hours scored better.")
print("3. Internal marks are higher for students who passed.")
print("4. Attendance and study hours positively affect performance.")
print("5. Internal Marks are the strongest indicator of the final result.")

# 7. Conclusion
print("\nConclusion")
print("The analysis shows that Attendance, Study Hours, and Internal Marks sig
```

Education Analytics Dataset

	Student_ID	Attendance	Study_Hours	Internal_Marks	Result
0	S01	92	5	45	Pass
1	S02	55	2	24	Fail
2	S03	78	4	38	Pass
3	S04	60	3	28	Fail
4	S05	88	5	42	Pass
5	S06	47	1	20	Fail
6	S07	95	6	48	Pass
7	S08	72	3	35	Pass
8	S09	50	2	22	Fail
9	S10	83	4	40	Pass
10	S11	66	3	30	Fail
11	S12	90	5	44	Pass
12	S13	58	2	25	Fail
13	S14	76	4	36	Pass
14	S15	62	3	29	Fail
15	S16	85	5	41	Pass
16	S17	45	1	18	Fail
17	S18	80	4	39	Pass
18	S19	68	3	32	Pass
19	S20	53	2	23	Fail

Rows and Columns: (20, 5)

Data Types

```
Student_ID      object
Attendance      int64
Study_Hours     int64
Internal_Marks  int64
Result          object
dtype: object
```

Missing Values

```
Student_ID      0
Attendance      0
Study_Hours     0
Internal_Marks  0
Result          0
dtype: int64
```

Duplicate Values: 0

Mean

```
Attendance      70.15
Study_Hours     3.35
Internal_Marks  32.95
dtype: float64
```


Median


```
Attendance      70.0
Study_Hours     3.0
Internal_Marks  33.5
dtype: float64
```

Minimum
Attendance 45
Study_Hours 1
Internal_Marks 18
dtype: int64

Maximum
Attendance 95
Study_Hours 6
Internal_Marks 48
dtype: int64

Standard Deviation
Attendance 15.938286
Study_Hours 1.424411
Internal_Marks 9.087151
dtype: float64





Important Patterns

- Students with higher attendance mostly passed.
- Students who studied more hours scored higher marks.
- Higher internal marks are associated with passing.
- Attendance and Internal Marks have a positive correlation.

Observations

1. Higher attendance improves the chances of passing.
2. Students studying 4-6 hours scored better.
3. Internal marks are higher for students who passed.
4. Attendance and study hours positively affect performance.
5. Internal Marks are the strongest indicator of the final result.

Conclusion

The analysis shows that Attendance, Study Hours, and Internal Marks significantly influence student performance. Students with higher attendance and study hours generally achieve better results.

The correlation matrix and pair plot are presented individually as they provide distinct analytical views. A pair plot inherently generates a grid of relationships, making it a comprehensive multi-plot figure on its own.

```
In [13]: # Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Create Dataset
data = {
    "Patient_ID": ["P01", "P02", "P03", "P04", "P05", "P06", "P07", "P08", "P09", "P10",
                  "P11", "P12", "P13", "P14", "P15", "P16", "P17", "P18", "P19", "P20"],

    "Age": [45, 32, 55, 28, 60, 40, 52, 35, 48, 30,
            62, 38, 50, 27, 58, 42, 46, 34, 53, 29],

    "Sugar_Level": [145, 98, 180, 90, 200, 120, 165, 105, 155, 95,
                   210, 115, 170, 88, 190, 130, 150, 100, 175, 92],

    "BP": [130, 118, 145, 110, 150, 125, 140, 120, 135, 115,
          155, 122, 142, 108, 148, 128, 132, 116, 144, 112],

    "BMI": [28, 23, 31, 22, 34, 26, 30, 24, 29, 23,
           35, 25, 30, 21, 33, 27, 28, 24, 31, 22],

    "Disease_Status": ["Yes", "No", "Yes", "No", "Yes", "No", "Yes", "No", "Yes", "No",
                      "Yes", "No", "Yes", "No", "Yes", "No", "Yes", "No", "Yes", "No"]
}

df = pd.DataFrame(data)

# Display Dataset
print("Healthcare Analytics Dataset")
print(df)

# 1. Number of Rows and Columns
print("\nRows and Columns:", df.shape)

# 2. Data Types
print("\nData Types")
print(df.dtypes)

# 3. Missing Values
print("\nMissing Values")
print(df.isnull().sum())

# Duplicate Values
print("\nDuplicate Values:", df.duplicated().sum())

# 4. Statistical Measures
print("\nMean")
print(df.mean(numeric_only=True))

print("\nMedian")
print(df.median(numeric_only=True))
```

```

print("\nMinimum")
print(df.min(numeric_only=True))

print("\nMaximum")
print(df.max(numeric_only=True))

print("\nStandard Deviation")
print(df.std(numeric_only=True))

# 5. Charts

# Histogram and Bar Chart
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.histplot(df["Sugar_Level"], bins=5, kde=True, color="skyblue", ax=axes[0])
axes[0].set_title("Sugar Level Distribution")
axes[0].set_xlabel("Sugar Level")
axes[0].set_ylabel("Frequency")
sns.countplot(x="Disease_Status", data=df, ax=axes[1])
axes[1].set_title("Disease Status Count")
plt.tight_layout()
plt.show()

# Scatter Plot and Box Plot
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.scatterplot(x="Sugar_Level", y="BMI", hue="Disease_Status", data=df, ax=axes[0])
axes[0].set_title("Sugar Level vs BMI")
sns.boxplot(x="Disease_Status", y="Sugar_Level", data=df, ax=axes[1])
axes[1].set_title("Sugar Level by Disease Status")
plt.tight_layout()
plt.show()

# 6. Important Patterns
print("\nImportant Patterns")
print("• Patients with higher sugar levels are mostly affected by disease.")
print("• Higher BMI is commonly observed in diseased patients.")
print("• Patients with high BP also tend to have disease.")
print("• Lower sugar levels are mostly associated with healthy patients.")

# 7. Observations
print("\nObservations")
print("1. High sugar level increases disease risk.")
print("2. Patients with higher BMI are more likely to have disease.")
print("3. High BP is common among diseased patients.")
print("4. Age has a moderate effect on disease.")
print("5. Sugar Level is the strongest indicator of disease.")

# 8. Conclusion
print("\nConclusion")
print("The EDA shows that Sugar Level, BMI, and BP are the major factors assoc

```

Healthcare Analytics Dataset

	Patient_ID	Age	Sugar_Level	BP	BMI	Disease_Status
0	P01	45	145	130	28	Yes
1	P02	32	98	118	23	No
2	P03	55	180	145	31	Yes
3	P04	28	90	110	22	No
4	P05	60	200	150	34	Yes
5	P06	40	120	125	26	No
6	P07	52	165	140	30	Yes
7	P08	35	105	120	24	No
8	P09	48	155	135	29	Yes
9	P10	30	95	115	23	No
10	P11	62	210	155	35	Yes
11	P12	38	115	122	25	No
12	P13	50	170	142	30	Yes
13	P14	27	88	108	21	No
14	P15	58	190	148	33	Yes
15	P16	42	130	128	27	No
16	P17	46	150	132	28	Yes
17	P18	34	100	116	24	No
18	P19	53	175	144	31	Yes
19	P20	29	92	112	22	No

Rows and Columns: (20, 6)

Data Types

```

Patient_ID      object
Age             int64
Sugar_Level     int64
BP              int64
BMI             int64
Disease_Status  object
dtype: object

```

Missing Values

```

Patient_ID      0
Age             0
Sugar_Level     0
BP              0
BMI             0
Disease_Status  0
dtype: int64

```

Duplicate Values: 0

Mean

```

Age             43.20
Sugar_Level     138.65
BP              129.75
BMI             27.30
dtype: float64

```

Median

```

Age             43.5



```


```
Sugar_Level    137.5
BP             129.0
BMI            27.5
dtype: float64
```

```
Minimum
Age      27
Sugar_Level  88
BP      108
BMI      21
dtype: int64
```

```
Maximum
Age      62
Sugar_Level  210
BP      155
BMI      35
dtype: int64
```

```
Standard Deviation
Age      11.316499
Sugar_Level  40.334490
BP      14.523574
BMI      4.256265
dtype: float64
```





Important Patterns

- Patients with higher sugar levels are mostly affected by disease.
- Higher BMI is commonly observed in diseased patients.
- Patients with high BP also tend to have disease.
- Lower sugar levels are mostly associated with healthy patients.

Observations

1. High sugar level increases disease risk.
2. Patients with higher BMI are more likely to have disease.
3. High BP is common among diseased patients.
4. Age has a moderate effect on disease.
5. Sugar Level is the strongest indicator of disease.

Conclusion

The EDA shows that Sugar Level, BMI, and BP are the major factors associated with disease status. Patients with higher values are more likely to have the disease.

```
In [14]: # Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Create Dataset
data = {
    "Product_ID": ["PR01", "PR02", "PR03", "PR04", "PR05", "PR06", "PR07", "PR08", "PR09", "PR10", "PR11", "PR12", "PR13", "PR14", "PR15", "PR16", "PR17", "PR18", "PR19", "PR20"],
    "Category": ["Grocery", "Snacks", "Dairy", "Fruits", "Vegetables", "Grocery", "Snacks", "Dairy", "Fruits", "Vegetables", "Grocery", "Snacks", "Dairy", "Fruits", "Vegetables", "Grocery", "Snacks", "Dairy", "Fruits", "Vegetables"],
    "Price": [50, 30, 45, 60, 40, 80, 25, 55, 70, 35, 90, 20, 65, 75, 45, 100, 35, 50, 80, 30],
    "Quantity_Sold": [20, 35, 18, 25, 30, 15, 40, 16, 22, 28, 12, 45, 14, 20, 26, 10, 32, 19, 18, 35],
}
```

```

        "Discount": [5, 3, 2, 5, 4, 6, 3, 2, 5, 4,
                     6, 2, 3, 5, 4, 7, 3, 2, 6, 3],

        "Revenue": [1000, 1050, 810, 1500, 1200, 1200, 1000, 880, 1540, 980,
                    1080, 900, 910, 1500, 1170, 1000, 1120, 950, 1440, 1050]
    }

    # Create DataFrame
    df = pd.DataFrame(data)

    # Display Dataset
    print("Retail Sales Dataset")
    print(df)

    # 1. Number of Rows and Columns
    print("\nRows and Columns:", df.shape)

    # 2. Check Data Types
    print("\nData Types")
    print(df.dtypes)

    # 3. Missing Values
    print("\nMissing Values")
    print(df.isnull().sum())

    # Duplicate Values
    print("\nDuplicate Values:", df.duplicated().sum())

    # 4. Statistical Measures
    print("\nMean")
    print(df.mean(numeric_only=True))

    print("\nMedian")
    print(df.median(numeric_only=True))

    print("\nMinimum")
    print(df.min(numeric_only=True))

    print("\nMaximum")
    print(df.max(numeric_only=True))

    print("\nStandard Deviation")
    print(df.std(numeric_only=True))

    # ----- Charts -----

    # 1. Histogram and Bar Chart
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))
    sns.histplot(df["Revenue"], bins=5, kde=True, ax=axes[0])
    axes[0].set_title("Revenue Distribution")

    category_revenue = df.groupby("Category")["Revenue"].sum()

```

```

category_revenue.plot(kind="bar", color="skyblue", ax=axes[1])
axes[1].set_title("Revenue by Category")
axes[1].set_xlabel("Category")
axes[1].set_ylabel("Revenue")
plt.tight_layout()
plt.show()

# 3. Scatter Plot and Box Plot
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.scatterplot(x="Price", y="Revenue", hue="Category", data=df, ax=axes[0])
axes[0].set_title("Price vs Revenue")
sns.boxplot(x="Category", y="Revenue", data=df, ax=axes[1])
axes[1].set_title("Revenue by Category")
plt.tight_layout()
plt.show()

# 5. Important Patterns
print("\nImportant Patterns")
print("• Fruits category has the highest revenue.")
print("• Snacks category has the highest quantity sold.")
print("• Higher-priced products generally earn more revenue.")
print("• Discount has less impact on revenue than price.")

# 6. Observations
print("\nObservations")
print("1. Fruits category generated the highest revenue.")
print("2. Snacks products recorded the highest sales.")
print("3. Revenue increases with product price.")
print("4. Discount has only a small effect on revenue.")
print("5. Price and Quantity Sold are the main factors affecting revenue.")

# 7. Conclusion
print("\nConclusion")
print("The EDA shows that Price and Quantity Sold significantly influence Reve

```

Retail Sales Dataset

	Product_ID	Category	Price	Quantity_Sold	Discount	Revenue
0	PR01	Grocery	50	20	5	1000
1	PR02	Snacks	30	35	3	1050
2	PR03	Dairy	45	18	2	810
3	PR04	Fruits	60	25	5	1500
4	PR05	Vegetables	40	30	4	1200
5	PR06	Grocery	80	15	6	1200
6	PR07	Snacks	25	40	3	1000
7	PR08	Dairy	55	16	2	880
8	PR09	Fruits	70	22	5	1540
9	PR10	Vegetables	35	28	4	980
10	PR11	Grocery	90	12	6	1080
11	PR12	Snacks	20	45	2	900
12	PR13	Dairy	65	14	3	910
13	PR14	Fruits	75	20	5	1500
14	PR15	Vegetables	45	26	4	1170
15	PR16	Grocery	100	10	7	1000
16	PR17	Snacks	35	32	3	1120
17	PR18	Dairy	50	19	2	950
18	PR19	Fruits	80	18	6	1440
19	PR20	Vegetables	30	35	3	1050

Rows and Columns: (20, 6)

Data Types

```
Product_ID      object
Category        object
Price           int64
Quantity_Sold   int64
Discount        int64
Revenue         int64
dtype: object
```

Missing Values

```
Product_ID      0
Category        0
Price           0
Quantity_Sold   0
Discount        0
Revenue         0
dtype: int64
```

Duplicate Values: 0

Mean

```
Price           54.0
Quantity_Sold   24.0
Discount        4.0
Revenue         1114.0
dtype: float64
```

Median





```
Price           50.0
```

Quantity_Sold 21.0
Discount 4.0
Revenue 1050.0
dtype: float64

Minimum
Price 20
Quantity_Sold 10
Discount 2
Revenue 810
dtype: int64

Maximum
Price 100
Quantity_Sold 45
Discount 7
Revenue 1540
dtype: int64

Standard Deviation
Price 22.803509
Quantity_Sold 9.673621
Discount 1.555973
Revenue 221.416493
dtype: float64



Important Patterns

- Fruits category has the highest revenue.
- Snacks category has the highest quantity sold.
- Higher-priced products generally earn more revenue.
- Discount has less impact on revenue than price.

Observations

1. Fruits category generated the highest revenue.
2. Snacks products recorded the highest sales.
3. Revenue increases with product price.
4. Discount has only a small effect on revenue.
5. Price and Quantity Sold are the main factors affecting revenue.

Conclusion

The EDA shows that Price and Quantity Sold significantly influence Revenue. Fruits are the most profitable category, while Snacks have the highest sales volume.

```
In [15]: # Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Create Dataset
data = {
    "Customer_ID": ["C01", "C02", "C03", "C04", "C05", "C06", "C07", "C08", "C09", "C10",
                    "C11", "C12", "C13", "C14", "C15", "C16", "C17", "C18", "C19", "C20"],

    "Income": [45000, 25000, 60000, 22000, 50000, 30000, 70000, 28000, 55000, 24000,
               65000, 35000, 48000, 26000, 75000, 32000, 52000, 27000, 68000, 23000],

    "Credit_Score": [720, 580, 750, 560, 710, 600, 780, 590, 730, 570,
                     760, 620, 700, 585, 790, 610, 725, 575, 770, 555],

    "Loan_Amount": [200000, 150000, 300000, 120000, 250000, 180000, 350000, 160000, 270
                    320000, 200000, 240000, 140000, 400000, 190000, 260000, 150000, 330

    "Employment_Type": ["Salaried", "Self-employed", "Salaried", "Unemployed", "Sal
                        "Self-employed", "Salaried", "Self-employed", "Salaried", "
                        "Salaried", "Self-employed", "Salaried", "Unemployed", "Sal
                        "Self-employed", "Salaried", "Self-employed", "Salaried", "

    "Loan_Status": ["Approved", "Rejected", "Approved", "Rejected", "Approved",
                    "Rejected", "Approved", "Rejected", "Approved", "Rejected",
                    "Approved", "Rejected", "Approved", "Rejected", "Approved",
                    "Rejected", "Approved", "Rejected", "Approved", "Rejected"]
}

# Create DataFrame
df = pd.DataFrame(data)

# Display Dataset
print("Banking Loan Approval Dataset")
print(df)
```

```

# 1. Number of Rows and Columns
print("\nRows and Columns:", df.shape)

# 2. Data Types
print("\nData Types")
print(df.dtypes)

# 3. Missing Values
print("\nMissing Values")
print(df.isnull().sum())

# Duplicate Values
print("\nDuplicate Values:", df.duplicated().sum())

# 4. Statistical Measures
print("\nMean")
print(df.mean(numeric_only=True))

print("\nMedian")
print(df.median(numeric_only=True))

print("\nMinimum")
print(df.min(numeric_only=True))

print("\nMaximum")
print(df.max(numeric_only=True))

print("\nStandard Deviation")
print(df.std(numeric_only=True))

# ----- Charts -----

# 1. Histogram and Bar Chart
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.histplot(df["Credit_Score"], bins=5, kde=True, ax=axes[0])
axes[0].set_title("Credit Score Distribution")

loan_count = df["Loan_Status"].value_counts()
loan_count.plot(kind="bar", color=["green", "red"], ax=axes[1])
axes[1].set_title("Loan Approval Status")
axes[1].set_xlabel("Loan Status")
axes[1].set_ylabel("Count")
plt.tight_layout()
plt.show()

# 3. Scatter Plot and Box Plot
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.scatterplot(x="Income", y="Credit_Score", hue="Loan_Status", data=df, ax=axes[0])
axes[0].set_title("Income vs Credit Score")
sns.boxplot(x="Loan_Status", y="Income", data=df, ax=axes[1])
axes[1].set_title("Income by Loan Status")
plt.tight_layout()

```

```
plt.show()

# 5. Important Patterns
print("\nImportant Patterns")
print("• Customers with higher income are mostly approved.")
print("• Higher credit scores increase loan approval chances.")
print("• Approved customers generally have higher income.")
print("• Loan approval mainly depends on income and credit score.")

# 6. Observations
print("\nObservations")
print("1. Customers with high income are more likely to get loans.")
print("2. Higher credit scores lead to higher approval rates.")
print("3. Rejected customers usually have lower income.")
print("4. Income and Credit Score have a positive relationship.")
print("5. Credit Score is an important factor in loan approval.")

# 7. Conclusion
print("\nConclusion")
print("The analysis shows that Income and Credit Score play a major role in lo
```


Banking Loan Approval Dataset

	Customer_ID	Income	Credit_Score	Loan_Amount	Employment_Type	Loan_Status
0	C01	45000	720	200000	Salaried	Approved
1	C02	25000	580	150000	Self-employed	Rejected
2	C03	60000	750	300000	Salaried	Approved
3	C04	22000	560	120000	Unemployed	Rejected
4	C05	50000	710	250000	Salaried	Approved
5	C06	30000	600	180000	Self-employed	Rejected
6	C07	70000	780	350000	Salaried	Approved
7	C08	28000	590	160000	Self-employed	Rejected
8	C09	55000	730	270000	Salaried	Approved
9	C10	24000	570	130000	Unemployed	Rejected
10	C11	65000	760	320000	Salaried	Approved
11	C12	35000	620	200000	Self-employed	Rejected
12	C13	48000	700	240000	Salaried	Approved
13	C14	26000	585	140000	Unemployed	Rejected
14	C15	75000	790	400000	Salaried	Approved
15	C16	32000	610	190000	Self-employed	Rejected
16	C17	52000	725	260000	Salaried	Approved
17	C18	27000	575	150000	Self-employed	Rejected
18	C19	68000	770	330000	Salaried	Approved
19	C20	23000	555	125000	Unemployed	Rejected

Rows and Columns: (20, 6)

Data Types

```
Customer_ID      object
Income           int64
Credit_Score     int64
Loan_Amount      int64
Employment_Type  object
Loan_Status      object dtype:
object
```

Missing Values

```
Customer_ID      0
Income           0
Credit_Score     0
Loan_Amount      0
Employment_Type  0
Loan_Status      0
dtype: int64
```

Duplicate Values: 0

Mean

```
Income           43000.0
Credit_Score     664.0
Loan_Amount      223250.0
dtype: float64
```

Median





```
Income           40000.0
Credit_Score     660.0
```

Loan_Amount 200000.0
dtype: float64

Minimum
Income 22000
Credit_Score 555
Loan_Amount 120000
dtype: int64

Maximum
Income 75000
Credit_Score 790
Loan_Amount 400000
dtype: int64

Standard Deviation
Income 17923.815383
Credit_Score 85.526235
Loan_Amount 83733.207272
dtype: float64



Important Patterns

- Customers with higher income are mostly approved.
- Higher credit scores increase loan approval chances.
- Approved customers generally have higher income.
- Loan approval mainly depends on income and credit score.

Observations

1. Customers with high income are more likely to get loans.
2. Higher credit scores lead to higher approval rates.
3. Rejected customers usually have lower income.
4. Income and Credit Score have a positive relationship.
5. Credit Score is an important factor in loan approval.

Conclusion

The analysis shows that Income and Credit Score play a major role in loan approval. Customers with higher values have a greater chance of getting their loans approved.

```
In [16]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
data = {
    "Farm_ID": ["F01", "F02", "F03", "F04", "F05", "F06", "F07", "F08", "F09", "F10",
               "F11", "F12", "F13", "F14", "F15", "F16", "F17", "F18", "F19", "F20"],
    "Rainfall_mm": [850, 600, 900, 500, 780, 650, 920, 720, 550, 800,
                   880, 620, 760, 940, 580, 820, 700, 960, 540, 790],
    "Temperature": [28, 32, 27, 34, 29, 31, 26, 30, 33, 28,
                   27, 32, 29, 26, 33, 28, 30, 25, 34, 29],
    "Fertilizer_kg": [60, 45, 65, 40, 55, 48, 70, 52, 42, 58,
                    62, 46, 54, 72, 44, 60, 50, 75, 41, 56],
    "Soil_Type": ["Loamy", "Sandy", "Loamy", "Sandy", "Clay",
                  "Sandy", "Loamy", "Clay", "Sandy", "Loamy",
                  "Loamy", "Sandy", "Clay", "Loamy", "Sandy",
                  "Clay", "Clay", "Loamy", "Sandy", "Clay"],
    "Crop_Yield_kg": [3200, 2100, 3500, 1800, 3000, 2300, 3700, 2800, 1900, 3100,
                    3400, 2200, 2950, 3800, 2000, 3150, 2700, 3900, 1850, 3050]
}

# Create DataFrame
df = pd.DataFrame(data)

# Display Dataset
print("Agriculture Crop Yield Dataset")
print(df)

# 1. Number of Rows and Columns
print("\nRows and Columns:", df.shape)

# 2. Data Types
```

```

print("\nData Types")
print(df.dtypes)

# 3. Missing Values
print("\nMissing Values")
print(df.isnull().sum())

# Duplicate Values
print("\nDuplicate Values:", df.duplicated().sum())

# 4. Statistical Measures
print("\nMean")
print(df.mean(numeric_only=True))

print("\nMedian")
print(df.median(numeric_only=True))

print("\nMinimum")
print(df.min(numeric_only=True))

print("\nMaximum")
print(df.max(numeric_only=True))

print("\nStandard Deviation")
print(df.std(numeric_only=True))

# ----- Charts -----

# 1. Histogram and Bar Chart
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.histplot(df["Crop_Yield_kg"], bins=5, kde=True, ax=axes[0])
axes[0].set_title("Crop Yield Distribution")
axes[0].set_xlabel("Crop Yield (kg)")
axes[0].set_ylabel("Frequency")

soil_yield = df.groupby("Soil_Type")["Crop_Yield_kg"].mean()
soil_yield.plot(kind="bar", color="green", ax=axes[1])
axes[1].set_title("Average Crop Yield by Soil Type")
axes[1].set_xlabel("Soil Type")
axes[1].set_ylabel("Average Crop Yield (kg)")
plt.tight_layout()
plt.show()

# 3. Scatter Plot and Box Plot
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.scatterplot(x="Rainfall_mm", y="Crop_Yield_kg", hue="Soil_Type", data=df,
axes[0].set_title("Rainfall vs Crop Yield")
sns.boxplot(x="Soil_Type", y="Crop_Yield_kg", data=df, ax=axes[1])
axes[1].set_title("Crop Yield by Soil Type")
plt.tight_layout()
plt.show()

print("\nImportant Patterns")

```

```
print("• Higher rainfall generally results in higher crop yield.")
print("• Loamy soil produces better crop yield than Sandy soil.")
print("• Higher fertilizer usage is associated with higher crop production.")
print("• Temperature has a moderate impact on crop yield.")

print("\nObservations")
print("1. Farms with higher rainfall produce higher crop yield.")
print("2. Loamy soil has the highest average crop yield.")
print("3. Sandy soil gives comparatively lower crop yield.")
print("4. Fertilizer usage positively influences crop production.")
print("5. Rainfall and fertilizer are the major factors affecting crop yield.")
print("\nConclusion")
print("The analysis shows that rainfall, soil type, and fertilizer usage signi
```

Agriculture Crop Yield Dataset

	Farm_ID	Rainfall_mm	Temperature	Fertilizer_kg	Soil_Type	Crop_Yield_kg
0	F01	850	28	60	Loamy	3200
1	F02	600	32	45	Sandy	2100
2	F03	900	27	65	Loamy	3500
3	F04	500	34	40	Sandy	1800
4	F05	780	29	55	Clay	3000
5	F06	650	31	48	Sandy	2300
6	F07	920	26	70	Loamy	3700
7	F08	720	30	52	Clay	2800
8	F09	550	33	42	Sandy	1900
9	F10	800	28	58	Loamy	3100
10	F11	880	27	62	Loamy	3400
11	F12	620	32	46	Sandy	2200
12	F13	760	29	54	Clay	2950
13	F14	940	26	72	Loamy	3800
14	F15	580	33	44	Sandy	2000
15	F16	820	28	60	Clay	3150
16	F17	700	30	50	Clay	2700
17	F18	960	25	75	Loamy	3900
18	F19	540	34	41	Sandy	1850
19	F20	790	29	56	Clay	3050

Rows and Columns: (20, 6)

Data Types

```
Farm_ID          object
Rainfall_mm      int64
Temperature      int64
Fertilizer_kg    int64
Soil_Type        object
Crop_Yield_kg    int64
dtype: object
```

Missing Values

```
Farm_ID          0
Rainfall_mm      0
Temperature      0
Fertilizer_kg    0
Soil_Type        0
Crop_Yield_kg    0
dtype: int64
```

Duplicate Values: 0

Mean

```
Rainfall_mm      743.00
Temperature      29.55
Fertilizer_kg    54.75
Crop_Yield_kg    2820.00
dtype: float64
```

Median



```
Rainfall_mm      770.0
```


Temperature 29.0
Fertilizer_kg 54.5
Crop_Yield_kg 2975.0
dtype: float64

Minimum
Rainfall_mm 500
Temperature 25
Fertilizer_kg 40
Crop_Yield_kg 1800
dtype: int64

Maximum
Rainfall_mm 960
Temperature 34
Fertilizer_kg 75
Crop_Yield_kg 3900
dtype: int64

Standard Deviation
Rainfall_mm 144.335136
Temperature 2.762055
Fertilizer_kg 10.507516
Crop_Yield_kg 681.407213
dtype: float64





Important Patterns

- Higher rainfall generally results in higher crop yield.
- Loamy soil produces better crop yield than Sandy soil.
- Higher fertilizer usage is associated with higher crop production.
- Temperature has a moderate impact on crop yield.

Observations

1. Farms with higher rainfall produce higher crop yield.
2. Loamy soil has the highest average crop yield.
3. Sandy soil gives comparatively lower crop yield.
4. Fertilizer usage positively influences crop production.
5. Rainfall and fertilizer are the major factors affecting crop yield.

Conclusion

The analysis shows that rainfall, soil type, and fertilizer usage significantly influence crop yield. Loamy soil with adequate rainfall and fertilizer gives the best crop production.